

Implementing an efficient and scalable service-oriented architecture for the Apertium machine translation platform

Pasquale Minervini

Dipartimento di Informatica
Università degli Studi di Bari
Via E. Orabona 4, 70125 Bari, Italy
p.minervini@gmail.com

Abstract

Service Oriented Architecture (SOA) is a paradigm for organising and using distributed services that may be under the control of different ownership domains and implemented using various technology stacks. In some contexts, an organisation using an IT infrastructure implementing the SOA paradigm can take a great benefit from the integration, in its business processes, of efficient machine translation (MT) services to overcome language barriers. This paper describes the architecture and the design patterns used to develop an MT service that is efficient, scalable and easy to integrate in new and existing business processes. The service is based on Apertium, a free/open-source rule-based machine translation platform.

1 Introduction

Service Oriented Architecture is an architectural paradigm providing a set of principles of governing concepts used during phases of systems development and integration. In such an architecture, functionalities are packaged as interoperable, loosely coupled services that may be used to build infrastructures enabling those with needs (consumers) and those with capabilities (providers) to interact across different domains of technology and ownership.

Several new trends in the computer industry rely upon SOA as their enabling foundation, including the automation of Business Process Management (BPM) and the multitude of new architecture and design patterns generally referred to as Web 2.0 (O'Reilly, 2005).

In some contexts, an organisation using an IT infrastructure implementing the SOA paradigm can take a great benefit from the integration, in its business processes, of an efficient machine translation service to overcome language barriers; for instance, it could be integrated in collaborative environments where people, who have no language in common, attempt to communicate with each other, or in knowledge extraction processes, where data is not available in a language that can be understood by the domain experts or the knowledge extraction tools being used.

A machine translation service was created using Apertium¹ (Armentano-Oller et al., 2005), a free/open-source rule-based machine translation platform for its translation capabilities, and on libTextCat², a library implementing n-gram based text categorisation (Cavnar and Trenkle, 1994) for language recognition, which provides an inexpensive and highly effective way of recognising the language used in documents; this result is reached by using small-sized (about 10K) fingerprints of the desired languages rather than resorting to more complicated and costly methods such as natural language parsing or assembling detailed lexicons. The libTextCat library is also

¹<http://www.apertium.org/>

²<http://software.wise-guys.nl/libtextcat/>

used by Bitextor (Esplà-Gomis, 2009), a system to harvest translation memories from multilingual websites.

Efficiency and scalability are critical for the service since, especially in collaborative environments, it should be able to sustain a heavy load of traffic. In this paper, the techniques and design patterns used to implement the machine translation service will be described and it will be compared to the standalone system.

2 Service APIs

Our service provides the two following capabilities:

- Translation – for automatic translation of free text from a source language to a destination language;
- Language recognition – for automatic language guessing of free text;

In SOA, interoperability between services is achieved by using standard languages for the description of service interfaces and the communications among services. A widely accepted technique for implementing SOA consists in making use of Web Services (Erl, 2005); a Web Service is defined by the W3C as “a software system designed to support interoperable machine-to-machine interaction over a network. It has an interface described in a machine-processable format (specifically WSDL). Other systems interact with the Web service in a manner prescribed by its description using SOAP-messages, typically conveyed using HTTP with an XML serialization in conjunction with other Web-related standards.” (Brown and Haas, 2004).

Alternative standards to SOAP are XML-RPC (Winer, 1999), a remote procedure call protocol which uses XML to encode its calls and HTTP as a transport mechanism, and Representational State Transfer (REST) (Fielding, 2000), a style of software architecture for distributed hypermedia systems such as the World Wide Web.

Our service provides XML-RPC, SOAP and REST interfaces to the translation and language recognition functionalities. All the interfaces follow the schema outlined in tables 1 and 2 to expose, respectively, the translation and the lan-

parameters	text
	source language
	destination language
returns	translation
	detected source language

Table 1: Parameters and return value(s) for the Translate method.

parameters	text
returns	detected language

Table 2: Parameters and return value(s) for the Detect method.

guage detection functionalities; those can be subsumed by the following methods:

- Translate: receives three parameters called `text`, `source language` and `destination language` containing, respectively the text to be translated, the source language and the destination language, and returns a `translation` value containing the translated text; if the source language is omitted, then language recognition is used to guess it, and the guessed language is returned in the `detected source language` value.
- Detect: receives three parameters called `text` containing free text, and returns a `detected language` value containing the language used by the text.

In addition, our service provides a `language pairs` method that returns a sequence of all the language pairs supported by the translation system, each represented by a pair containing the corresponding `source language` and the `destination language`.

In all methods, languages are represented by their ISO 639-1 (ISO:639-1, 2002) code.

3 Internal architecture of the service

Apertium is a transfer-based machine translation system which uses finite-state transducers for lexical processing, hidden Markov models (HMMs) for part-of-speech tagging and finite-state-based chunking for structural transfer. Its translation en-

Relying on a resource pool is designed to result in the following improvements for our rule-based

machine translation service:

Performance – By preventing repetitious acquisition and release of resources;

Stability – Because repetitious acquisition and release of resource can increase the risk of system instability (for example, repetitious acquisition and release of memory can lead to fragmentation problems);

Scalability – As resources can be used also in different types of translation tasks, avoiding the allocation of a complete set of resources for each different translation task (for example, translation tasks on different language pairs using different dictionaries can make use of the same resource for managing the removal and restoration of formatting).

4 Results

To evaluate the efficiency of our service, which we will refer to as `apertium-service`, we compared the time it requires to compute and answer to a translation request from Spanish to English with the `apertium-en-es` language pair³ with the time required by the following systems:

- `apertium`, a console application implemented as a part of the Apertium project;




Figure 2: Comparison in the “Sentence Length - Time” space between `apertium` and `apertium-service`; measurements are in *string length* for the Sentence Length dimension and in *ms* for the Time dimension.

³PUT THE VERSION OR SVN REVISION HERE




Figure 3: Comparison in the “Number of Threads- Time” space between `apertium` and `apertium-service`.

All the experiments were run on a server with four 2GHz Dual-Core AMD Opteron processors and 4GB of main memory, using the GNU/Linux operating system. Both `apertium-service` and `moses-service` were accepting translation requests in XML-RPC format, and the free text used for timing all the systems was also taken from Europarl corpus. Figure ?? shows the time required to translate increasingly longer sentences for all systems (values in the time dimension are shown on a logarithmic scale), and figure 2 only for `apertium-service` and `apertium`.

Scalability for `apertium` and `apertium-service` have been evaluated by calculating the average time required by the two systems to answer to 1,024 translations requests sequentially sent by a variable number of clients; the requests consisted to translating the longest sentence from the Europarl evaluation corpus (679 characters) from Spanish to English. Figure 3 shows the results of this comparison.

5 Future work

UMLS concept identification in non-English medical documents: MetaMap (Aronson, 2001) is an application that allows mapping text to UMLS Metathesaurus⁴ concepts, which have proved to be useful for many applications,

⁴The UMLS[®] Metathesaurus[®] (Schuyler et al., 1993) provides a representation of biomedical knowledge consisting of concepts classified by semantic type and both hierarchical and non-hierarchical relationships among the concepts.

including decision support systems, management of patient records, information retrieval and data mining within the biomedical domain.

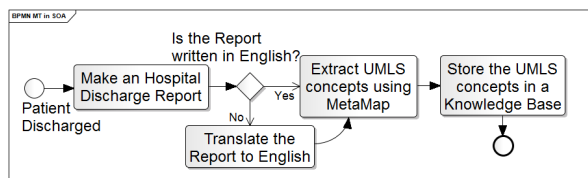


Figure 4: Representation of a business process in which a clinical document, if written in a language different than English, is first translated to English and then processed using MetaMap to extract UMLS concepts.

Currently MetaMap is only available for English free text, which makes it difficult the use of UMLS Metathesaurus to represent concepts from biomedical texts written in languages other than English. A possible way to overcome this limitation consists in using machine translation specific for the biomedical domain to translate the free text from its original language to English, and then process it as in Figure 4. This approach is discussed in Carrero et al. (2008).

6 Conclusions

We presented *apertium-service*, a machine translation service based on Apertium, a free/open-source rule-based machine translation platform. It has been shown to be competitive compared to the baseline system in both efficiency and scalability.

Source code for our service is released under the GNU General Public Licence version 3⁵ and is available on the Apertium SVN repository.⁶

Acknowledgements

Development for this project was funded as part of the Google Summer of Code⁷ programme. Many thanks go to Jimmy O'Regan, Francis Tyers and others involved in the Apertium Project, for their constant help. Additionally I am grateful to the anonymous reviewers for their invaluable comments and suggestions on an earlier version of this paper.

⁵<http://www.gnu.org/licenses/gpl.html>

⁶<http://apertium.svn.sourceforge.net/svnroot/apertium/trunk/apertium-service>

⁷<http://code.google.com/soc/>

References

- Armentano-Oller, C., Corbí-Bellot, A. M., Forcada, M. L., Ginestí-Rosell, M., Bonev, B., Ortiz-Rojas, S., Pérez-Ortiz, J. A., Ramírez-Sánchez, G., and Sánchez-Martínez, F. (2005). An open-source shallow-transfer machine translation toolbox: consequences of its release and availability. In *OSMaTran: Open-Source Machine Translation, A workshop at Machine Translation Summit X*, pages 23–30.
- Aronson, A. R. (2001). Effective mapping of biomedical text to the umls metathesaurus: the metamap program. *Proc AMIA Symp*, pages 17–21.
- Brown, A. and Haas, H. (2004). Web services glossary. World Wide Web Consortium, Note NOTE-ws-gloss-20040211.
- Carrero, F. M., Cortizo, J. C., Gómez, J. M., and de Buenaga, M. (2008). In the development of a spanish metamap. In *CIKM '08: Proceeding of the 17th ACM conference on Information and knowledge management*, pages 1465–1466, New York, NY, USA. ACM.
- Cavnar, W. B. and Trenkle, J. M. (1994). N-gram-based text categorization. In *In Proceedings of SDAIR-94, 3rd Annual Symposium on Document Analysis and Information Retrieval*, pages 161–175.
- Erl, T. (2005). *Service-Oriented Architecture : Concepts, Technology, and Design*. Prentice Hall PTR.
- Esplà-Gomis, M. (2009). Bitextor: a Free/Open-source Software to Harvest Translation Memories from Multilingual Websites. In *Proceedings of MT Summit XII*, Ottawa, Canada. Association for Machine Translation in the Americas.
- Fielding, R. T. (2000). *Architectural Styles and the Design of Network-based Software Architectures*. PhD thesis, University of California, Irvine.
- ISO:639-1 (2002). Iso 639-1:2002 – codes for the representation of names of languages – part 1: Alpha-2 code.

- Kircher, M. and Jain, P. (2004). *Pattern-Oriented Software Architecture Volume 3: Patterns for Resource Management*. Wiley.
- O'Reilly, T. (2005). *What Is Web 2.0: Design Patterns and Business Models for the Next Generation of Software*.
- Schuyler, P. L., Hole, W. T., Tuttle, M. S., and Sherertz, D. D. (1993). The umls metathesaurus: representing different views of biomedical concepts. *Bull Med Libr Assoc*, 81(2):217–222.
- Winer, D. (1999). XML/RPC specification. Technical report, Userland Software.